

Human ratings are highly reliable where individual responses exist

291 of 335 tasks have item-level individual ratings (median $r = 0.90$)

overall median $r = 0.90$

