

Human ratings are highly reliable

item-level split-half, 291 of 335 tasks with individual responses · overall median $r = 0.90$

